

Name :Dinesh Kumar 

Roll : 225229108

Lab-5 : Stemming and Lemmatization on Movie Dataset

```
In [18]: from zipfile import ZipFile
import glob
import pandas as pd
import nltk
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import linear_kernel
from nltk.corpus import stopwords
import warnings
warnings.filterwarnings('ignore')
```

EXERCISE-1

```
In [19]: file_name = "movies.zip"
with ZipFile(file_name, 'r') as zip:
    zip.printdir()
```

File Name	Modified	Size
movies/	2018-01-19 08:32:38	0
movies/12 Angry Men.txt	2018-01-17 20:40:42	1007
movies/12 Years a Slave.txt	2018-01-17 20:42:50	6451
movies/4 Months, 3 Weeks and 2 Days.txt	2018-01-17 20:37:10	1151
movies/All About Eve.txt	2018-01-17 20:33:18	1346
movies/American Graffiti.txt	2018-01-17 20:44:30	3417
movies/Boys n the Hood.txt	2018-01-17 20:27:14	1970
movies/Casablanca.txt	2018-01-17 20:26:26	1896
movies/Citizen Kane.txt	2018-01-17 20:23:56	1483
movies/Gone with the Wind.txt	2018-01-17 20:38:10	1318
movies/Hoop Dreams.txt	2018-01-17 20:34:12	7909
movies/Manchester by the Sea.txt	2018-01-17 20:40:06	3674
movies/Moonlight.txt	2018-01-17 20:31:42	2323
movies/My Left Foot.txt	2018-01-17 20:38:50	1115
movies/Pan's Labyrinth.txt	2018-01-17 20:32:18	4431
movies/Psycho.txt	2018-01-17 20:34:46	3727
movies/Ran.txt	2018-01-17 20:43:48	2207
movies/Singin' in the Rain.txt	2018-01-17 20:29:42	782
movies/Some Like It Hot.txt	2018-01-17 20:35:40	7489
movies/The Godfather.txt	2018-01-17 20:25:32	4293
movies/Three Colors Red.txt	2018-01-17 20:28:22	2892

```
In [20]: nltk.download('punkt')
nltk.download('stopwords')
stop_words = set(stopwords.words('english'))

[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\1mscdsa20\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\1mscdsa20\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
In [21]: from nltk.stem import PorterStemmer
ps = PorterStemmer()
tokenizer = nltk.tokenize.WhitespaceTokenizer()
from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()
from nltk.stem import LancasterStemmer
ls = LancasterStemmer()
```

```
In [24]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readlines()
        print(contents)
        print("*****")
        print(" ")
```

["Lumet's origins as a director of teledrama may well be obvious here in his first film, but there is no denying the suitability of his style - sweaty close-ups, gritty monochrome 'realism', one-set claustrophobia - to his subject. Scripted by Reginald Rose from his own teleplay, the story is pretty contrived - during a murder trial, one man's doubts about the accused's guilt gradually overcome the rather less-than-democratic prejudices of the other eleven members of the jury - but the treatment is tense, lucid, and admirably economical. Fonda, though typecast as the bastion of liberalism, gives a nicely underplayed performance, while Cobb, Marshall and Begley in particular are highly effective in support. But what really transforms the piece from a rather talky demonstration that a man is innocent until proven guilty, is the consistently taut, sweltering atmosphere, created largely by Boris Kaufman's excellent camerawork. The result, however devoid of action, is a strangely realistic thriller."]

['There are movies to which the critical response lags far behind the emotional one. Two days after seeing 12 Years a Slave, British director Steve McQueen's adaptation of the 1853 memoir of a free black man kidnapped into slavery, the still awaiting delivery of the narrative that would permit me to analyze

A. How many sentences in each file?

```
In [25]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readlines()
        for row in contents:
            sent_text = nltk.sent_tokenize(row)
            print("sentence tokenize ", len(sent_text))
```

```
sentence tokenize 5
sentence tokenize 4
sentence tokenize 0
sentence tokenize 4
sentence tokenize 0
sentence tokenize 5
sentence tokenize 0
sentence tokenize 4
sentence tokenize 0
sentence tokenize 5
sentence tokenize 0
sentence tokenize 7
sentence tokenize 0
sentence tokenize 2
sentence tokenize 4
sentence tokenize 0
sentence tokenize 2
sentence tokenize 0
sentence tokenize 2
```

B. How many tokens in each file?

```
In [29]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readlines()
        for row1 in contents:
            words = nltk.word_tokenize(row1)
            print("word tokenize ", len(words))
```

```
word tokenize 181
word tokenize 119
word tokenize 20
word tokenize 276
word tokenize 9
word tokenize 70
word tokenize 49
word tokenize 98
word tokenize 242
word tokenize 67
word tokenize 131
word tokenize 157
word tokenize 69
word tokenize 66
word tokenize 39
word tokenize 25
word tokenize 50
word tokenize 208
word tokenize 100
word tokenize 569
```

C. How many tokens excluding stop words in each file?

```
In [30]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readlines()
        filtered_sentence = [w for w in words if not w in stop_words]
        print("stopwords ", len(filtered_sentence))
```

```
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
```

D. How many unique stems (ie., stemming) in each file? (Use PorterStemmer)

```
In [31]: def port_stemSentence(sentence):
tokenizer = nltk.tokenize.WhitespaceTokenizer()
tok = tokenizer.tokenize(sentence)
filtered_sentence = [w for w in tok if not w in stop_words]
stem_sentence = []
for word in filtered_sentence:
    stem_sentence.append(ps.stem(word))
return len(stem_sentence)
```

```
In [32]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readline()
        print("porter_stemming ")
        print(porter_stemSentence(contents))
```

```
porter_stemming
96
porter_stemming
83
porter_stemming
20
porter_stemming
138
porter_stemming
63
porter_stemming
64
porter_stemming
20
porter_stemming
51
porter_stemming
131
porter_stemming
27
porter_stemming
53
porter_stemming
87
porter_stemming
35
porter_stemming
93
porter_stemming
23
porter_stemming
34
porter_stemming
52
porter_stemming
38
porter_stemming
33
porter_stemming
282
```

E. How many unique stems (ie., stemming) in each file? (Use LancasterStemmer)

```
In [33]: def lan_stemSentence(sentence):  
    tokenizer = nltk.tokenize.WhitespaceTokenizer()  
    tok = tokenizer.tokenize(sentence)  
    filtered_sentence = [w for w in tok if not w in stop_words]  
    stem_sentence = []  
    for word in filtered_sentence:  
        stem_sentence.append(ls.stem(word))  
    return len(stem_sentence)
```

```
In [34]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readline()
        print("lancaster_stemming ")
        print(port_stemSentence(contents))
```

```
lancaster_stemming
96
lancaster_stemming
83
lancaster_stemming
20
lancaster_stemming
138
lancaster_stemming
63
lancaster_stemming
64
lancaster_stemming
20
lancaster_stemming
51
lancaster_stemming
131
lancaster_stemming
27
lancaster_stemming
53
lancaster_stemming
87
lancaster_stemming
35
lancaster_stemming
93
lancaster_stemming
23
lancaster_stemming
34
lancaster_stemming
52
lancaster_stemming
38
lancaster_stemming
33
lancaster_stemming
282
```

F. How many unique words (ie., lemmatization) in each file? (Use WordNetLemmatizer)


```
In [53]: import nltk  
nltk.download('wordnet')
```

```
[nltk_data] Downloading package wordnet to  
[nltk_data] C:\Users\1mscdsa20\AppData\Roaming\nltk_data...  
[nltk_data] Package wordnet is already up-to-date!
```

Out[53]: True

```
In [54]: def lemmSentence(sentence):  
    tokenizer = nltk.tokenize.WhitespaceTokenizer()  
    tok = tokenizer.tokenize(sentence)  
    filtered_sentence = [w for w in tok if not w in stop_words]  
    lemm_sentence = []  
    for word in filtered_sentence:  
        lemm_sentence.append(lemmatizer.lemmatize(word))  
    return len(lemm_sentence)
```

```
In [55]: for file in files:
          with open(file, 'r', encoding='cp1252') as f:
              contents = f.readline()
              print("lemmatization ")
              print(lemmSentence(contents))
```

```
lemmatization
96
lemmatization
83
lemmatization
20
lemmatization
138
lemmatization
63
lemmatization
64
lemmatization
20
lemmatization
51
lemmatization
131
lemmatization
27
lemmatization
53
lemmatization
87
lemmatization
35
lemmatization
93
lemmatization
23
lemmatization
34
lemmatization
52
lemmatization
38
lemmatization
33
lemmatization
282
```

EXERCISE-2

Step-1 For each movie:

Tokenize terms and build list of tokens

```
In [56]: tok = []
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.read()
        let=tokenizer.tokenize(contents)
        tok.append(let)
tok
```

```
Out[56]: [['Lumet's',
'origins',
'as',
'a',
'director',
'of',
'teledrama',
'may',
'well',
'be',
'obvious',
'here',
'in',
'his',
'first',
'film,',
'but',
'there',
'is',
',',
',']
```

Find lemmatized words from the tokens

```
In [57]: import nltk
nltk.download('omw-1.4')
```

```
[nltk_data] Downloading package omw-1.4 to
[nltk_data] C:\Users\1mscdsa20\AppData\Roaming\nltk_data...
[nltk_data] Package omw-1.4 is already up-to-date!
```

```
Out[57]: True
```

```
In [58]: tok_lem = []
         for i in tok:
             for j in i:
                 to_lem = lemmatizer.lemmatize(j)
                 tok_lem.append(to_lem)
         tok_lem
```

```
Out[58]: ["Lumet's",
          'origin',
          'a',
          'a',
          'director',
          'of',
          'teledrama',
          'may',
          'well',
          'be',
          'obvious',
          'here',
          'in',
          'his',
          'first',
          'film,',
          'but',
          'there',
          'is',
          ',']
```

Step-2

Build Term-Document matrix using TfidfVectorizer

```
In [59]: for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.read()
        tok = tokenizer.tokenize(contents)
        filtered_sentence = [w for w in tok if not w in stop_words]
        tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
        features = tfidf.fit_transform(filtered_sentence)
        df = pd.DataFrame(features.todense(), columns=tfidf.get_feature_names())
        print(df)
```

```
      man  one  rather
0  0.0  0.0    0.0
1  0.0  0.0    0.0
2  0.0  0.0    0.0
3  0.0  0.0    0.0
4  0.0  0.0    0.0
..  ...  ...    ...
91 0.0  0.0    0.0
92 0.0  0.0    0.0
93 0.0  0.0    0.0
94 0.0  0.0    0.0
95 0.0  0.0    0.0

[96 rows x 3 columns]
      12  all  almost  and  beautiful  black  but  children  comes  cotton  \
0  0.0  0.0    0.0  0.0    0.0    0.0  0.0    0.0    0.0    0.0
1  0.0  0.0    0.0  0.0    0.0    0.0  0.0    0.0    0.0    0.0
2  0.0  0.0    0.0  0.0    0.0    0.0  0.0    0.0    0.0    0.0
3  0.0  0.0    0.0  0.0    0.0    0.0  0.0    0.0    0.0    0.0
4  0.0  0.0    0.0  0.0    0.0    0.0  0.0    0.0    0.0    0.0
```

Step-3

Take vectors of any two movies and compute cosine similarity

```
In [60]: with open(files[5], 'r', encoding='cp1252') as f:
          contents = f.read()
          tok = tokenizer.tokenize(contents)
          filtered_sentence = [w for w in tok if not w in stop_words]
          tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
          movie1 = tfidf.fit_transform(filtered_sentence)
          print(movie1)
```

```
(1, 10)      1.0
(5, 2)       1.0
(12, 13)     1.0
(15, 5)      1.0
(18, 10)     1.0
(31, 20)     1.0
(35, 12)     1.0
(37, 3)      1.0
(38, 9)      1.0
(45, 10)     1.0
(46, 11)     1.0
(48, 19)     1.0
(49, 16)     1.0
(53, 8)      1.0
(54, 4)      1.0
(56, 19)     1.0
(62, 20)     1.0
(65, 12)     1.0
(69, 7)      1.0
(72, 18)     0.5773502691896258
(72, 14)     0.5773502691896258
(72, 17)     0.5773502691896258
(77, 6)      1.0
(78, 18)     0.5773502691896258
(78, 14)     0.5773502691896258
:           :
(108, 7)     1.0
(118, 5)     1.0
(121, 13)    1.0
(124, 12)    1.0
(128, 6)     1.0
(134, 10)    1.0
(138, 15)    1.0
(143, 15)    1.0
(148, 7)     1.0
(152, 1)     1.0
(154, 1)     1.0
(156, 1)     1.0
(165, 9)     1.0
(166, 0)     1.0
(172, 4)     1.0
(173, 2)     1.0
(174, 8)     1.0
(177, 10)    1.0
(179, 3)     1.0
(180, 0)     1.0
(188, 20)    1.0
(193, 7)     1.0
(194, 11)    1.0
```

(196, 12)	1.0
(203, 10)	1.0

```
In [61]: with open(files[10], 'r', encoding='cp1252') as f:
          contents = f.read()
          tok = tokenizer.tokenize(contents)
          filtered_sentence = [w for w in tok if not w in stop_words]
          tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
          movie2 = tfidf.fit_transform(filtered_sentence)
          print(movie2)
```

```
(0, 15)      1.0
(1, 27)      1.0
(2, 34)      1.0
(3, 6)       1.0
(4, 8)       1.0
(7, 26)      1.0
(11, 22)     1.0
(13, 19)     1.0
(15, 20)     1.0
(17, 0)      1.0
(29, 11)     1.0
(34, 16)     1.0
(46, 35)     1.0
(52, 43)     1.0
(53, 20)     1.0
(62, 11)     1.0
(66, 20)     1.0
(67, 10)     1.0
(71, 14)     1.0
(73, 2)      1.0
(74, 18)     1.0
(77, 37)     1.0
(78, 12)     1.0
(81, 39)     1.0
(82, 20)     1.0
:           :
(323, 34)    1.0
(324, 25)    1.0
(331, 42)    1.0
(332, 19)    1.0
(333, 40)    1.0
(336, 23)    1.0
(337, 29)    1.0
(342, 31)    1.0
(343, 33)    1.0
(345, 38)    1.0
(353, 3)     1.0
(354, 11)    1.0
(356, 24)    1.0
(359, 28)    1.0
(361, 27)    1.0
(362, 34)    1.0
(366, 43)    1.0
(369, 22)    1.0
(371, 30)    1.0
(373, 41)    1.0
(379, 4)     1.0
(381, 36)    1.0
(383, 7)     1.0
```



```
(384, 39)    1.0  
(385, 4)    1.0
```

```
In [62]: doc1 = movie1[0:10]  
doc2 = movie1[:]  
score = linear_kernel(doc1,doc2)  
print(score)
```

```
[[0. 0. 0. ... 0. 0. 0.]  
 [0. 1. 0. ... 0. 0. 0.]  
 [0. 0. 0. ... 0. 0. 0.]  
 ...  
 [0. 0. 0. ... 0. 0. 0.]  
 [0. 0. 0. ... 0. 0. 0.]  
 [0. 0. 0. ... 0. 0. 0.]]
```

```
In [ ]:
```